

# Computerlinguistik II

Vorlesung im SoSe 2019  
(M-GSW-10)

**Prof. Dr. Udo Hahn**

Lehrstuhl für Computerlinguistik  
Institut für Germanistische Sprachwissenschaft  
Friedrich-Schiller-Universität Jena

<http://www.julielab.de>

# Two Paradigms for NLP

- Symbolic Specification Paradigm
  - Manual acquisition procedures
  - Lab-internal activities
  - Intuition and (few!) subjectively generated examples drive progress based on individual (competence) judgments
    - “I have a system that parses all of my nine-teen sentences!”

# Symbolic Specification Paradigm

- **Manual rule specification**
  - Source: linguist's intuition
- **Manual lexicon specification**
  - Source: linguist's intuition
- **Each lab has its own (home-grown) set of NLP software**
  - Hampers reusability
  - Limits scientific progress
  - Waste of human and monetary resources (we “burnt” thousands of Ph.D. student all over the world ☹)

# Shortcomings of the “Classical” Linguistic Approach

- Huge amounts of background knowledge req.
  - Lexicons (approx. 100,000 – 150,000 entries)
  - Grammars (>> 15,000 – 20,000 rules)
  - Semantics (>> 15,000 – 20,000 rules)
- As the linguistic and conceptual coverage of classical linguistic systems increases (slowly), it still remains insufficient; systems also reveal ‘spurious’ ambiguity, and, hence, tend to become overly “brittle” and unmaintainable
- More fail-soft behavior is required at the expense of ... ? (e.g., full-depth understanding)

# Two Paradigms for NLP

## • Symbolic Specification Paradigm

- Manual acquisition procedures
- Lab-internal activities
- Intuition and (few!) subjectively generated examples drive progress based on individual (competence) judgments
  - “I have a system that parses all of my nine-teen sentences!”





## • Empirical (Learning) Paradigm

- Automatic acquisition procedures
- Community-wide sharing of common knowledge and resources
- Large and ‘representative’ data sets drive progress according to experimental standards
  - “The system was tested on 1,7 million words taken from the WSJ segment of the MUC-7 data set and produced 4.9% parsing errors, thus yielding a statistically significant 1.6% improvement over the best result by parser X on the same data set & a 40.3% improvement over the baseline system!”

# Empirical Paradigm


- Large repositories of language data
  - Corpora (plain or annotated, i.e., enriched by meta-data)
- Large, community-wide shared repositories of language processing modules
  - Tokenizers, POS taggers, chunkers, NE recognizers, ...
- Shared repositories of machine learning algos
- Automatic acquisition of linguistic knowledge
  - Applying ML algos to train linguistic processors by using large corpora with valid linguistic metadata (linguist as educated data supplier, „language expert“) rather than manual intuition (linguist as creative rule inventor)
- Shallow analysis rather than deep understanding
- Large, community-wide self-managed, task-oriented competitions, comparative evaluation rounds
- Change of mathematics:
  - Statistics rather than algebra and logics

# Paradigm Shift – We Exchanged our Textbooks...



# POS Tagging

A severe infection ended the pregnancy .



DET ADJ NOUN VERB DET NOUN ST



# Penn Treebank Tag Set

In total,  
45 tags

Tag	Description	Examples
.	sentence terminator	. ! ?
DT	determiner	all an many such that the them these this
JJ	adjective, numeral	first oiled separable battery-powered
NN	common noun	cabbage thermostat investment
PRP	personal pronoun	herself him it me one oneself theirs they
IN	preposition	among out within behind into next
VB	verb (base form)	ask assess assign begin break bring
VBD	verb (past tense)	asked assessed assigned began broke
WP	WH-pronoun	that what which who whom

# Transformation Rules for Tagging [Brill, 1995]

- Initial State: Based on a number of features, guess the most likely POS tag for a given word:
  - die/DET Frau/NOUN ,/COMMA die/DET singt/VFIN
- Learn transformation rules to reduce errors:
  - *Change DET to PREL whenever the preceding word is tagged as COMMA*
- Apply learned transformation rules:
  - die/DET Frau/NOUN, /COMMA die/PREL singt/VFIN

# First 20 Transformation Rules

#	Change Tag		Condition
	From	To	
1	NN	VB	Previous tag is <i>TO</i>
2	VBP	VB	One of the previous three tags is <i>MD</i>
3	NN	VB	One of the previous two tags is <i>MD</i>
4	VB	NN	One of the previous two tags is <i>DT</i>
5	VBD	VBN	One of the previous three tags is <i>VBZ</i>
6	VBN	VBD	Previous tag is <i>PRP</i>
7	VBN	VBD	Previous tag is <i>NNP</i>
8	VBD	VBN	Previous tag is <i>VBD</i>
9	VBP	VB	Previous tag is <i>TO</i>
10	POS	VBZ	Previous tag is <i>PRP</i>
11	VB	VBP	Previous tag is <i>NNS</i>
12	VBD	VBN	One of previous three tags is <i>VBP</i>
13	IN	WDT	One of next two tags is <i>VB</i>
14	VBD	VBN	One of previous two tags is <i>VB</i>
15	VB	VBP	Previous tag is <i>PRP</i>
16	IN	WDT	Next tag is <i>VBZ</i>
17	IN	DT	Next tag is <i>NN</i>
18	JJ	NNP	Next tag is <i>NNP</i>
19	IN	WDT	Next tag is <i>VBD</i>
20	JJR	RBR	Next tag is <i>JJ</i>

Taken from: Brill (1995), Transformation-Based Error-Driven Learning

# Towards Statistical Models of Natural Language Processing ...

# Letter-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
-

# Letter-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
- **W**

# Letter-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
- **Wh**

# Letter-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
- **Wha**



# Letter-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
- **What**

# Letter-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
- **What d**

# Letter-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
- **What do**

# Letter-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
- **What do you think the next letter is?**

# Word-based Language Models

- **Shannon's Game**
- **Guess the next letter:**
  - **What do you think the next letter is?**
- **Guess the next word:**
-

# Word-based Language Models

- **Shannon's Game**
- Guess the next letter:
  - What do you think the next letter is?
- Guess the next word:
  - We

# Word-based Language Models

- **Shannon's Game**
- Guess the next letter:
  - What do you think the next letter is?
- Guess the next word:
  - We are

# Word-based Language Models

- **Shannon's Game**
- Guess the next letter:
  - What do you think the next letter is?
- Guess the next word:
  - We are now



# Word-based Language Models

- **Shannon's Game**
- Guess the next letter:
  - What do you think the next letter is?
- Guess the next word:
  - We are now entering

# Word-based Language Models

- **Shannon's Game**
- Guess the next letter:
  - What do you think the next letter is?
- Guess the next word:
  - We are now entering statistical

# Word-based Language Models

- **Shannon's Game**
- Guess the next letter:
  - What do you think the next letter is?
- Guess the next word:
  - We are now entering statistical territory

# Approximating Natural Language Words

- **zero-order approximation:**  
letter sequences are independent of  
each other and all equally probable:
  - xfoml rxkhrjffjuj zlpwcwkcy  
ffjeyvkcqsghyd

# Approximating Natural Language Words

- **first-order approximation:**  
letters are independent, but occur  
with the frequencies of English text:
  - ocro hli rgwr nmielwis eu ll  
nbnesebya th eei alhenhtppa oobttva  
nah

# Approximating Natural Language Words

- **second-order approximation:**  
the probability that a letter appears depends on the previous letter
  - on ie antsoutinys are t inctore st bes  
deamy achin d ilonasive tucoowe at  
teasonare fuzo tizin andy tobe seace  
ctisbe

# Approximating Natural Language Words


- **third-order approximation:**  
the probability that a certain letter appears depends on the two previous letters
  - in no ist lat whey cratict froure birs  
grocid pondenome of demonstures  
of the reptagin is regoactiona of cre

# Approximating Natural Language Words

- Higher frequency trigrams for different languages:
  - English: THE, ING, ENT, ION
  - German: EIN, ICH, DEN, DER
  - French: ENT, QUE, LES, ION
  - Italian: CHE, ERE, ZIO, DEL
  - Spanish: QUE, EST, ARA, ADO



# Zipfsches Gesetz



Wortverteilung im Vergleich zu einer einfachen Zipf-Verteilung ( $\sim 1/n$ . Wortanzahl: 70;  
Texte aus: <http://www.gutenberg.org/dirs/etext04/8effi10.txt>)

# Terminology

- **Sentence:** unit of written language
- **Utterance:** unit of spoken language
- **Word Form:** the inflected form that appears literally in the corpus
- **Lemma:** lexical forms having the same stem, part of speech, and word sense
- **Types (V):** number of distinct words that might appear in a corpus (vocabulary size)
- **Tokens ( $N_T$ ):** total number of words in a corpus (note:  $V \ll N_T$ )
- **Types seen so far (T):** number of distinct words seen so far in corpus (note:  $T \leq V \ll N_T$ )

# Word-based Language Models

- A model that enables one to compute the probability, or likelihood, of a sentence  $S$ ,  $P(S)$ .
- Simple: Every word follows every other word with equal probability (0-gram)
  - Assume  $|V|$  is the size of the vocabulary  $V$
  - Likelihood of sentence  $S$  of length  $n$  is  $1/|V| \times 1/|V| \dots \times 1/|V|$
  - If English has 100,000 words, the probability of each next word is  $1/100000 = .00001$

# Relative Frequency vs. Conditional Probability

- Smarter: *Relative* Frequency

Probability of each next word is related to word frequency within a corpus (unigram)

- Likelihood of sentence  $S = P(w_1) \times P(w_2) \times \dots \times P(w_n)$
- Assumes probability of each word is independent of probabilities of other words

# Relative Frequency vs. Conditional Probability

- Smarter: *Relative* Frequency

Probability of each next word is related to word frequency within a corpus (unigram)

- Likelihood of sentence  $S = P(w_1) \times P(w_2) \times \dots \times P(w_n)$
- Assumes probability of each word is independent of probabilities of other words

- Even smarter: *Conditional* Probability

Look at probability given previous words (n-gram)

- Likelihood of sentence  $S = P(w_1) \times P(w_2|w_1) \times \dots \times P(w_n|w_{n-1})$
- Assumes probability of each word is dependent on probabilities of previous words

# Generalization of Conditional Probability via Chain Rule

- Conditional Probability for Two Events,  $A_1$  and  $A_2$ 
  - $P(A_1, A_2) = P(A_1) \cdot P(A_2|A_1)$
- **Chain Rule** generalizes to multiple ( $n$ ) events
  - $P(A_1, \dots, A_n) =$   
$$P(A_1) \times P(A_2|A_1) \times P(A_3|A_1, A_2) \times \dots \times P(A_n|A_1 \dots A_{n-1})$$
  - Examples:
    - $P(\text{the dog}) = P(\text{the}) \times P(\text{dog} | \text{the})$
    - $P(\text{the dog bites}) = P(\text{the}) \times P(\text{dog} | \text{the}) \times P(\text{bites} | \text{the dog})$

# Relative Frequencies and Conditional Probabilities

- Relative word frequencies are better than equal probabilities for all words
  - In a corpus with 10K word types, each word would have  $P(w) = 1/10K$
  - Does not match our intuitions that different words are more likely to occur
    - (e.g. “the” vs. “shop” vs. “aardvark”)
- Conditional probability is more useful than individual relative word frequencies
  - **dog** may be relatively rare in a corpus
  - but if we see **barking**,  $P(\text{dog}|\text{barking})$  may be large

# Probability for a Word String

- In general, the probability of a complete string of words  $w_1^n = w_1 \dots w_n$  is

$$\begin{aligned} P(w_1^n) \\ &= P(w_1)P(w_2/w_1)P(w_3/w_1 \ w_2) \dots P(w_n/w_1 \dots w_{n-1}) \\ &= \prod_{k=1}^n P(w_k | w_1^{k-1}) \end{aligned}$$

- But this approach to determining the probability of a word sequence gets to be computationally very expensive and suffers from sparse data



# Markov Assumption (basic idea)

- How do we (efficiently) compute  $P(w_n | w_1^{n-1})$ ?
- Trick (!): Instead of  $P(\text{rabbit} | \text{I saw } \underline{a})$ , we use  $P(\text{rabbit} | \underline{a})$ .
  - This lets us collect statistics in practice via a bigram model:  $P(\text{the barking dog}) = P(\text{the} | \text{<start>}) \times P(\text{barking} | \text{the}) \times P(\text{dog} | \text{barking})$

# Markov Assumption (the very idea)

- Markov models are the class of probabilistic language models that assume that we can predict the probability of some future unit *without looking too far* into the past
  - Specifically, for  $N=2$  (bigram):
    - $P(w_1^n) \approx \prod_{k=1}^n P(w_k | w_{k-1}); w_0 := \text{<start>}$
- Order of a Markov model: length of prior context
  - bigram is first order, trigram is second order, ...

# Statistical HMM-based Tagging

[Brants, 2000]

- *State transition probability*. Likelihood of a tag immediately following n other tags
    - $P_1(\text{Tag}_i \mid \text{Tag}_{i-1} \dots \text{Tag}_{i-n})$
  - *State emission probability*. Likelihood of a word given a tag
    - $P_2(\text{Word}_i \mid \text{Tag}_i)$
- 
- die/DET Frau/NOUN ,/COMMA die/DET or PREL singt/VFIN

# Trigrams for Tagging

- *State transition probabilities (trigrams):*

- $P_1(\text{DET} \mid \text{COMMA NOUN}) = 0.0007$

- $P_1(\text{PREL} \mid \text{COMMA NOUN}) = 0.01$

- *State emission probabilities:*

- $P_2(\text{die} \mid \text{DET}) = 0.7$

- $P_2(\text{die} \mid \text{PREL}) = 0.2$

- Compute probabilistic evidence for the tag being

- **DET:**  $P_1 \cdot P_2 = 0.0007 \cdot 0.7 = 0.00049$

- **PREL:**  $P_1 \cdot P_2 = 0.01 \cdot 0.2 = 0.002$

Taken from  
(POS-  
annotated)  
corpora

• die/DET Frau/NOUN ,/COMMA die/PREL singt/VFIN

# Inside (most) POS Taggers

- Lexicon look-up routines
- Morphological processing (not only deflection!)
- Unknown word handler, if lexicon look-up fails (based on statistical information)
- Ambiguity ranking (priority selection)

# Chunking

Arginine methylation of STAT1 modulates IFN induced transcription

# Chunking

[Arginine methylation] of [STAT1] modulates [IFN induced transcription]

# Shallow Parsing

[Arginine methylation of STAT1]<sub>NP</sub> [modulates]<sub>VP</sub> [IFN induced transcription]<sub>NP</sub>



# Shallow Parsing

[ [Arginine methylation]<sub>NP</sub> [of STAT1]<sub>PP</sub> ]<sub>NP</sub>

[Arginine methylation of STAT1]<sub>NP</sub> [modulates]<sub>VP</sub> [IFN induced transcription]<sub>NP</sub>

# Shallow Parsing

[ [IFN induced]<sub>AP</sub> [transcription]<sub>N</sub> ]<sub>NP</sub>

[ [Arginine methylation]<sub>NP</sub> [of STAT1]<sub>PP</sub> ]<sub>NP</sub>

[Arginine methylation of STAT1]<sub>NP</sub> [modulates]<sub>VP</sub> [IFN induced transcription]<sub>NP</sub>

# Deep Parsing

[ [IFN induced]<sub>AP</sub> [transcription]<sub>N</sub> ]<sub>NP</sub>

[ [ [Arginine]<sub>N</sub> [methylation]<sub>N</sub> ]<sub>NP</sub> [ [of]<sub>P</sub> [STAT1]<sub>N</sub> ]<sub>PP</sub> ]<sub>NP</sub>

[ [Arginine methylation]<sub>NP</sub> [of STAT1]<sub>PP</sub> ]<sub>NP</sub>

[Arginine methylation of STAT1]<sub>NP</sub> [ [modulates]<sub>V</sub> [IFN induced transcription]<sub>NP</sub> ]<sub>VP</sub>

# Deep Parsing

[ [ [IFN]<sub>N</sub> [induced]<sub>A</sub> ]<sub>AP</sub> [transcription]<sub>N</sub> ]<sub>NP</sub>

[ [IFN induced]<sub>AP</sub> [transcription]<sub>N</sub> ]<sub>NP</sub>

[ [ [Arginine]<sub>N</sub> [methylation]<sub>N</sub> ]<sub>NP</sub> [ [of]<sub>P</sub> [STAT1]<sub>N</sub> ]<sub>PP</sub> ]<sub>NP</sub>

[ [Arginine methylation]<sub>NP</sub> [of STAT1]<sub>PP</sub> ]<sub>NP</sub>

[Arginine methylation of STAT1]<sub>NP</sub> [ [modulates]<sub>V</sub> [IFN induced transcription]<sub>NP</sub> ]<sub>VP</sub>

# Chunking Principles

- Goal: divide a sentence into a sequence of chunks (ako phrases)
- Chunks are non-overlapping regions of a text
  - *[I] saw [a tall man] in [the park]*
- Chunks are non-exhaustive
  - not all words of a sentence are included in chunks
- Chunks are non-recursive
  - a chunk does not contain other chunks
- Chunks are mostly base NP chunks

[ *[the synthesis]*<sub>NP-base</sub> *of* *[long enhancer transcripts]*<sub>NP-base</sub> ]<sub>NP-complex</sub>

# The Shallow Syntax Pipeline


Tagging




Chunking




Parsing



BIO (or IOB)



Trees



# BIO Format for Base NPs

a	DT	B
mechanism	NN	I
that	WDT	B
increases	VBZ	O
NF-kappa	NN	B
B/I	NN	I
kappa	NN	I
B	NN	I
dissociation	NN	I
without	IN	O
affecting	VBG	O
the	DT	B
NF-kappa	NN	I
B	NN	I
translocation	NN	I
step	NN	I

# A Simple Chunking Technique

- Simple chunkers usually ignore lexical content
  - Only need to look at part-of-speech tags
- Basic steps in chunking
  - Chunking / Unchunking
  - Chinking
  - Merging / Splitting



# Regular Expression Basics

- “|” OR operator (explicit OR-ing)
  - “[a|e|i|o|u]” matches any occurrence of vowels
- “[abc]” matches any occurrence of either “a”, “b” or “c” (implicit OR-ing)
  - “gr[ae]y” matches “grey” or “gray” (but not “graey”)
- “.” matches arbitrary char
  - “d.g” matches “dag”, “dig”, “dog”, “dkg” ...
- “?” preceding expression/char may or may not occur
  - “colou?r” matches “colour” and “color”
- “+” preceding expression occurs at least one time
  - “(ab)+” matches “ab”, “abab”, “ababab”, ...
- “\*” preceding expression occurs null time or arbitrary often
  - “(ab)\*” matches “\_”, “ab”, “abab”, “ababab”, ...

# Chunking

- Define a regular expression that matches the sequences of tags in a chunk
  - `<DT>? <JJ>* <NN.??>`
- Chunk all matching subsequences
  - *A/DT red/JJ car/NN ran/VBD on/IN the/DT street/NN*
  - *[A/DT red/JJ car/NN] ran/VBD*  
*on/IN [the/DT street/NN]*
- If matching subsequences overlap, the first one gets priority
- **Unchunking** is the opposite of chunking

# Chinking

- A chink is a subsequence of the text that is not a chunk
- Define a regular expression that matches the sequences of tags in a chink
  - ( <VB.??> | <IN> )+
- Chunk anything that is not a matching subsequence
  - *A/DT red/JJ car/NN ran/VBD on/IN the/DT street/NN*
  - [*A/DT red/JJ car/NN*]  

*ran/VBD on/IN*

 [*the/DT street/NN*]  

chink

# Merging

- Combine adjacent chunks into a single chunk
- Define a regular expression that matches the sequences of tags on both sides of the point to be merged
  - Merge a chunk ending in “JJ” with a chunk starting with “NN”, i.e. left: <JJ>, right: <NN.>
- Chunk all matching subsequences
  - [A/DT *red*/JJ ] [ *car*/NN] *ran*/VBD  
*on*/IN *the*/DT *street*/NN
  - [A/DT *red*/JJ *car*/NN] *ran*/VBD  
*on*/IN *the*/DT *street*/NN
- **Splitting** is the opposite of merging

# Concluding Remarks

- Chunking – as the weakest form of syntactic structuring – relies on RegExs
- RegExs (formally) belong to the class of regular grammars
- Regular grammars and their (finite-state) automata have linear run-time complexity
- Standard CF grammars and their associated push-down automata have (at best) cubic run-time complexity
- Hence, there is a trade-off between different levels of richness of syntactic structures and gains/losses of run-time behavior

# What are Named Entities?

- Names of persons
  - *Dr. Jonathan Peeko, Professor Johnson*
- Names of companies or organizations
  - *Sony, United Nations, Texas Instruments, General Motors*
- Names of locations
  - *Paris, San Francisco, Rocky Mountains, Yellowstone Park*
- Date and time expressions
  - *Feb 17, 1973; 4.40p.m.; 16.40 Uhr; autumn 2000; last year*
- Addresses
  - *7 Ugly Way, Wolverhampton UH0 1Q5*
  - *udo.hahn@uni-jena.de*
- Names of proteins or genes or diseases,
  - *chloramphenicol acetyltransferase, NF-kappa B, SARS*
- Measure expressions
  - *420 kp, 21 l/m<sup>2</sup>, 37%, 900€*

# What are Named Entities?

- Names of persons
  - *Dr. Jonathan Peeko, Professor John*
- Names of companies
  - *Sony, Universal Motors*
- Names of locations
  - *Central Park*
- Dates
  - *1999*
- Addresses
  - *1010 105*
  - *udo.hahn@uni-jena.de*
- Names of proteins or genes or diseases,
  - *chloramphenicol acetyltransferase, NF-kappa B, SARS*
- Measure expressions
  - *420 kp, 21 l/m<sup>2</sup>, 37%, 900€*

named entities are  
intentionally excluded from  
the lexicon